

# System Programming Lab #8

---

2019-04-30

sp-tas

# Lab Assignment #4 : Kernel Lab

- Download skeleton code & pdf from eTL
  - kernellab-handout.tar, kernellab-handout.pdf
- Hand In #1 – Setup (Done)
  - capture your development environment
  - Upload your capture image eTL
    - 압축파일 양식 : [학번]\_[이름]\_kernellab\_setup.tar (or .zip etc) **(including below files)**
      - filename for part #1 : [학번]\_[이름]\_kernellab\_ptree.jpg (or .png, etc)
      - filename for part #2: [학번]\_[이름]\_kernellab\_paddr.jpg (or .png, etc)
- Hand In #2 – Your Implementation
  - Upload your files eTL
    - 압축파일 양식 : [학번]\_[이름]\_kernellab.tar (or .zip, etc)
    - Ex) 2017-12345\_홍길동\_kernellab.tar
  - A zip file should include
  - (1) a tarball of your implementation directory (2) report
    - tarball 양식 : kernellab-[학번].tar.gz eg) kernellab-2019-12345.tar.gz
    - Report 양식 : [학번]\_[이름]\_kernellab\_report.pdf (or .hwp, .txt etc)
- Please, **READ** the Hand-out and Lab material thoroughly!

# Lab Assignment #4 : Kernel Lab

- Step 1. Setup (Done)
  - (part #0) Load my own kernel module
- Step 2. Implementation
  - (part #1) Tracing process tree from process id
  - (part #2) Finding physical address using virtual address
- Assigned : April 23
- Deadline for Step 1. Setup : April 29, 23:59:59 (Delay NOT allowed)
- Deadline for Step 2. Implementation : May 13, 23:59:59
- Delay policy : Same as before
- Lab sessions will be
  - 4/23: Kernel lab part #0, #1
  - 4/30: Kernel lab part #2 ← TODAY
  - 5/7 : Kernel lab Q&A session

# Today's Lab

---

- Some useful tools in kernel programming
  - tmux, ctags, cscope
- Part#2. Finding physical address using virtual address
  - Remind translation of VA->PA in Computer Architecture
  - Assignment spec
  - How to begin
  - Testing your program
- Evaluation

# tmux (terminal multiplexer)

- Installation
  - `sudo apt-get install tmux`
- Basic Commands

Command	Description
<code>tmux</code>	start tmux
<code>tmux new -s &lt;name&gt;</code>	start tmux with <name>
<code>tmux ls</code>	shows the list of sessions
<code>tmux a #</code>	attach the detached-session
<code>tmux a -t &lt;name&gt;</code>	attach the detached-session to <name>
<code>tmux kill-session -t &lt;name&gt;</code>	kill the session <name>
<code>tmux kill-server</code>	kill the tmux server

# Split windows with tmux

- **1. ctrl + b (or a)** to type tmux command
- 2. Split vertically : `%(shift 5)`
- 3. Split horizontally : `”(shift ‘)`
- <https://tmuxguide.readthedocs.io/en/latest/tmux/tmux.html>

# Ctags

- What is Ctags?
  - A Tool that makes it easy to navigate big source code projects.
- Ctags generates **database** of tag file
  - for global variables, functions, macros, etc
  - to point where they are declared & defined
- Installation
  - `sudo apt-get install ctags (or exuberant-ctags)`
- Check
  - `ctags --version`
- Help
  - `Ctags --help`

# Ctags – how to make tags file

- Steps
  - 1. go to root directory of codes you want to navigate.
    - `cd /(where your root directory of code is)`
  - 2. generate tags file
    - type `ctags -R` (recursive)
    - or `ctags file1, file2, ...`
  - 3. Check tags file
    - `ls`

```
1. ta@sp3: ~/yschoi/malloclab/src (ssh)
ta@sp3:~/yschoi/malloclab/src$ ctags -R
ta@sp3:~/yschoi/malloclab/src$ ls
checkalign  clock.o  fcyc.o  ftimer.c  Makefile-handout  memlib.c  mm-explicit.c  mm.o  README-handout
checkalign.c  config.h  fsecs.c  ftimer.h  mdriver          memlib.h  mm.h          mm-test.c  tags
clock.c      fcyc.c   fsecs.h  ftimer.o  mdriver.c        memlib.o  mm-implicit.c  mm-tree.c
clock.h      fcyc.h   fsecs.o  Makefile  mdriver.o        mm.c      mm-naive.c    README
```

- 4. Remove tags file
  - `rm tags`



# Ctags – how to use

- Case 1. In code file
  - 1. place cursor on the keyword you want to locate where it is defined
  - 2. type **ctrl + ]**

```
227 if ((bp = mem_sbrk(size)) == (void *)-1)
228     return NULL;
229
230 /* Initialize free block header/footer and the epilogue header */
231 PUT(HDRP(bp), PACK(size, 0)); /* free block header */
232 PUT(FTRP(bp), PACK(size, 0)); /* free block footer */
233 PUT(HDRP(NEXT_BLKP(bp)), PACK(0, 1)); /* new epilogue header */
234
235 /* Coalesce if the previous block was free */
236 return coalesce(bp);
237 }
238 /* $end mmextendheap */
239
240 /*
241  * place - Place block of asize bytes at start of free block bp
242  *         and split if remainder would be at least minimum block size
243  */
244 /* $begin mmplace */
245 /* $begin mmplace-proto */
246 static void place(void *bp, size_t asize)
247 {
248     /* ... */
249 }
```

227:17 [62%]  
"mm.c" 369L, 9640C

```
58 void *mem_sbrk(int incr)
59 {
60     char *old_brk = mem_brk;
61
62     if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr)) {
63         errno = ENOMEM;
64         fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
65         return (void *)-1;
66     }
67     mem_brk += incr;
68     return (void *)old_brk;
69 }
70
71 /*
72  * mem_heap_lo - return address of the first heap byte
73  */
74 void *mem_heap_lo()
75 {
76     return (void *)mem_start_brk;
77 }
78
79 void *mem_heap_hi()
80 {
81     return (void *)mem_end_brk;
82 }
```

58:1 [61%]  
'memlib.c' 101L, 2270C

- 3. type **ctrl + t** to go back

# Ctags – how to use

- Case 2. In tags file
  - 1. **vi tags**
  - 2. type **:tj [tag name]** to find

```
34 FREE_MASK mm-implicit.c 65;" d file:
35 FTRP mm-implicit.c 71;" d file:
36 FTRP mm.c 71;" d file:
37 GET mm-implicit.c 62;" d file:
38 GET mm.c 62;" d file:
1:1 [Top]
:tj mem_sbrk
```

- 3. type **:po** to comeback

```
58 void *mem_sbrk(int incr)
59 {
60     char *old_brk = mem_brk;
61
62     if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr)) {
63         errno = ENOMEM;
64         fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
65         return (void *)-1;
66     }
67     mem_brk += incr;
68     return (void *)old_brk;
69 }
70
71 /*
72  * mem_heap_lo - return address of the first heap byte
73  */
74 void *mem_heap_lo()
75 {
76     return (void *)mem_start_brk;
77 }
58:1 [61%]
:po
```

# Where is linux kernel code?

- /usr/src/linux-headers-4.xx/

```
root@yschoi-VirtualBox:/usr/src/linux-headers-4.15.0-47# ls
arch      drivers  ipc      Makefile  scripts  tools
block     firmware Kbuild   mm         security ubuntu
certs     fs       Kconfig  net        sound    usr
crypto    include  kernel   new.c      spl      virt
Documentation  init    lib      samples   tags     zfs
```

```
root@yschoi-VirtualBox:/usr/src/linux-headers-4.15.0-47# vi tags
```

```
#include <linux/debugfs.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/uaccess.h>
#include <linux/list.h>
#include <linux/slab.h>

#define COMM_STR_SIZE 32
#define BLOB_SIZE 128

MODULE_LICENSE("GPL");

static struct dentry *dir, *inputdir, *ptreedir;
static struct task_struct *curr;

struct task_list {
    struct list_head list;
    pid_t pid;
    char comm[COMM_STR_SIZE];
};

struct task_list t_list;

struct debugfs_blob_wrapper p_tree;
char blob[BLOB_SIZE];

static void add_task(struct task_struct *task)
{
    struct task_list *node;

    node = (struct task_list*)kmalloc(sizeof(struct task_list), GFP_KERNEL);
    node->pid = task->pid;
    strncpy(node->comm, task->comm, COMM_STR_SIZE);

    list_add((struct list_head*)node, &(t_list.list));
}

static void build_blob(void)
{
    char buffer[BLOB_SIZE] = "";
    "dbfs_ptree.c" 104 lines, 2481 characters
```

# Using Ctags in Linux kernel code

```

!_TAG_FILE_FORMAT      2       /extended format; --format=1 will
not append ";" to lines/
!_TAG_FILE_SORTED      1       /0=unsorted, 1=sorted, 2=foldcase/
!_TAG_PROGRAM_AUTHOR    Darren Hiebert /dhiebert@users.sourceforge
e.net/
!_TAG_PROGRAM_NAME      Exuberant Ctags //
!_TAG_PROGRAM_URL      http://ctags.sourceforge.net /official
site/
!_TAG_PROGRAM_VERSION   5.9~svn20110310 //
$0      arch/mips/include/asm/mach-cavium-octeon/kernel-entry-init
.h      /^      dins      v0, $0, 0, 6$/"      v
$0      arch/mips/include/asm/mach-cavium-octeon/kernel-entry-init
.h      /^      sd      $0, -32768(v0)$/"      v
$0      arch/mips/include/asm/stackframe.h      /^
NG_S     $0, PT_R0(sp)$/"      v
$0      arch/mips/include/asm/stackframe.h      /^
r      v1, $0, v1$/"      v
$1      arch/mips/include/asm/asmmacro.h      /^      _cfcmsa $1
, MSA_CSR$/"      v
$1      arch/mips/include/asm/asmmacro.h      /^      sw      $1
, THREAD_MSA_CSR(\\thread)$/"      v
$10     arch/mips/include/asm/stackframe.h      /^      cf
i_ld     $10, PT_R10, \\docfi$/"      v
$10     arch/mips/include/asm/stackframe.h      /^      cf
i_st     $10, PT_R10, \\docfi$/"      v
$11     arch/mips/include/asm/stackframe.h      /^      cf
i_ld     $11, PT_R11, \\docfi$/"      v
$11     arch/mips/include/asm/stackframe.h      /^      cf
i_st     $11, PT_R11, \\docfi$/"      v
$12     arch/mips/include/asm/stackframe.h      /^      cf
i_ld     $12, PT_R12, \\docfi$/"      v
$12     arch/mips/include/asm/stackframe.h      /^      cf
i_st     $12, PT_R12, \\docfi$/"      v
$13     arch/mips/include/asm/stackframe.h      /^      cf
i_ld     $13, PT_R13, \\docfi$/"      v
$13     arch/mips/include/asm/stackframe.h      /^      cf
i_st     $13, PT_R13, \\docfi$/"      v
$14     arch/mips/include/asm/stackframe.h      /^      cf
i_ld     $14, PT_R14, \\docfi$/"      v
$14     arch/mips/include/asm/stackframe.h      /^      cf
i_st     $14, PT_R14, \\docfi$/"      v
:tj task_struct

```

```

#include <linux/debugfs.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/uaccess.h>
#include <linux/list.h>
#include <linux/slab.h>

#define COMM_STR_SIZE    32
#define BLOB_SIZE        128

MODULE_LICENSE("GPL");

static struct dentry *dir, *inputdir, *ptreedir;
static struct task_struct *curr;

LO
no
struct task_list {
    struct list_head list;
    pid_t pid;
    char comm[COMM_STR_SIZE];
};

cf struct task_list t_list;

cf struct debugfs_blob_wrapper p_tree;
char blob[BLOB_SIZE];

cf static void add_task(struct task_struct *task)
{
    struct task_list *node;

    node = (struct task_list*)kmalloc(sizeof(struct task_list), GFP_KERNEL);
    node->pid = task->pid;
    strncpy(node->comm, task->comm, COMM_STR_SIZE);

    list_add((struct list_head*)node, &(t_list.list));
}

cf static void build_blob(void)
{
    char buffer[BLOB_SIZE] = "";

    "dbfs_ptree.c" 104 lines, 2481 characters

```

# Using Ctags in Linux kernel code

```
        u8                exp_need_qs;

        /* Otherwise the compiler can store garbage here: */
        u8                pad;
    } b; /* Bits. */
    u32 s; /* Set of bits. */
};

enum perf_event_task_context {
    perf_invalid_context = -1,
    perf_hw_context = 0,
    perf_sw_context,
    perf_nr_task_contexts,
};

struct wake_q_node {
    struct wake_q_node *next;
};

struct task_struct {
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /*
     * For reasons of header soup (see current_thread_info()),
     * this
     * must be the first element of task_struct.
     */
    struct thread_info        thread_info;
#endif
    /* -1 unrunnable, 0 runnable, >0 stopped: */
    volatile long             state;

    /*
     * This begins the randomizable portion of task_struct. On
     * ly
     * scheduling-critical items should be added above here.
     */
    randomized_struct_fields_start

    void                    *stack;
    atomic_t                usage;
};

#include <linux/debugfs.h>
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/uaccess.h>
#include <linux/list.h>
#include <linux/slab.h>

#define COMM_STR_SIZE      32
#define BLOB_SIZE          128

MODULE_LICENSE("GPL");

static struct dentry *dir, *inputdir, *ptreedir;
static struct task_struct *curr;

struct task_list {
    struct list_head list;
    pid_t pid;
    char comm[COMM_STR_SIZE];
};

struct task_list t_list;

struct debugfs_blob_wrapper p_tree;
char blob[BLOB_SIZE];

static void add_task(struct task_struct *task)
{
    struct task_list *node;

    node = (struct task_list*)kmalloc(sizeof(struct task_list), GFP_KERNEL);
    node->pid = task->pid;
    strncpy(node->comm, task->comm, COMM_STR_SIZE);

    list_add((struct list_head*)node, &(t_list.list));
}

static void build_blob(void)
{
    char buffer[BLOB_SIZE] = "";

    "dbfs ptree.c" 104 lines, 2481 characters
}
```

# Cscope

- A tool to navigate in big source code.
- Diff with Ctags?
  - Able to locate **functions where they are called** too.
- Installation
  - `sudo apt-get install cscope`
- Check
  - `cscope --version`
- Help
  - `cscope --help`

```
ta@sp3:~/yschoi/malloclab/src$ cscope --version
cscope: version 15.8b
ta@sp3:~/yschoi/malloclab/src$
```

# Cscope – how to make cscope database file

- Steps
  - 1. go to root directory of codes you want to navigate.
    - `cd /(where your root directory of code is)`
  - 2. generate cscope database file
    - `find ./ -name '*[cCsShH]' > file_list`
    - `cscope -i file_list`
  - 3. Check cscope.out file
    - `ls`

```
ta@sp3:~/yschoi/malloclab/src$ find ./ -name '*[cCsShH]' > file_list
ta@sp3:~/yschoi/malloclab/src$ cscope -i file_list
ta@sp3:~/yschoi/malloclab/src$ ls
checkalign  cscope.out  fsecs.h    Makefile-handout  memlib.o    mm.o
checkalign.c  fcyc.c      fsecs.o    mdriver           mm.c         mm-test.c
clock.c       fcyc.h      ftimer.c   mdriver.c         mm-explicit.c mm-tree.c
clock.h       fcyc.o      ftimer.h   mdriver.o         mm.h         README
clock.o       file_list   ftimer.o   memlib.c          mm-implicit.c README-handout
config.h      fsecs.c     Makefile   memlib.h          mm-naive.c
```

- 4. Remove tags file
  - `rm cscope.out file_list`

# Cscope – how to use

- 1. type **cscope** to execute
- 2. type **ctrl+d** to break out

Cscope version 15.8b

Press the ? key for help

Find this C symbol:  
Find this global definition:  
Find functions called by this function:  
Find functions calling this function: **mem\_sbrk**  
Find this text string:  
Change this text string:  
Find this egrep pattern:  
Find this file:  
Find files #including this file:  
Find assignments to this symbol:

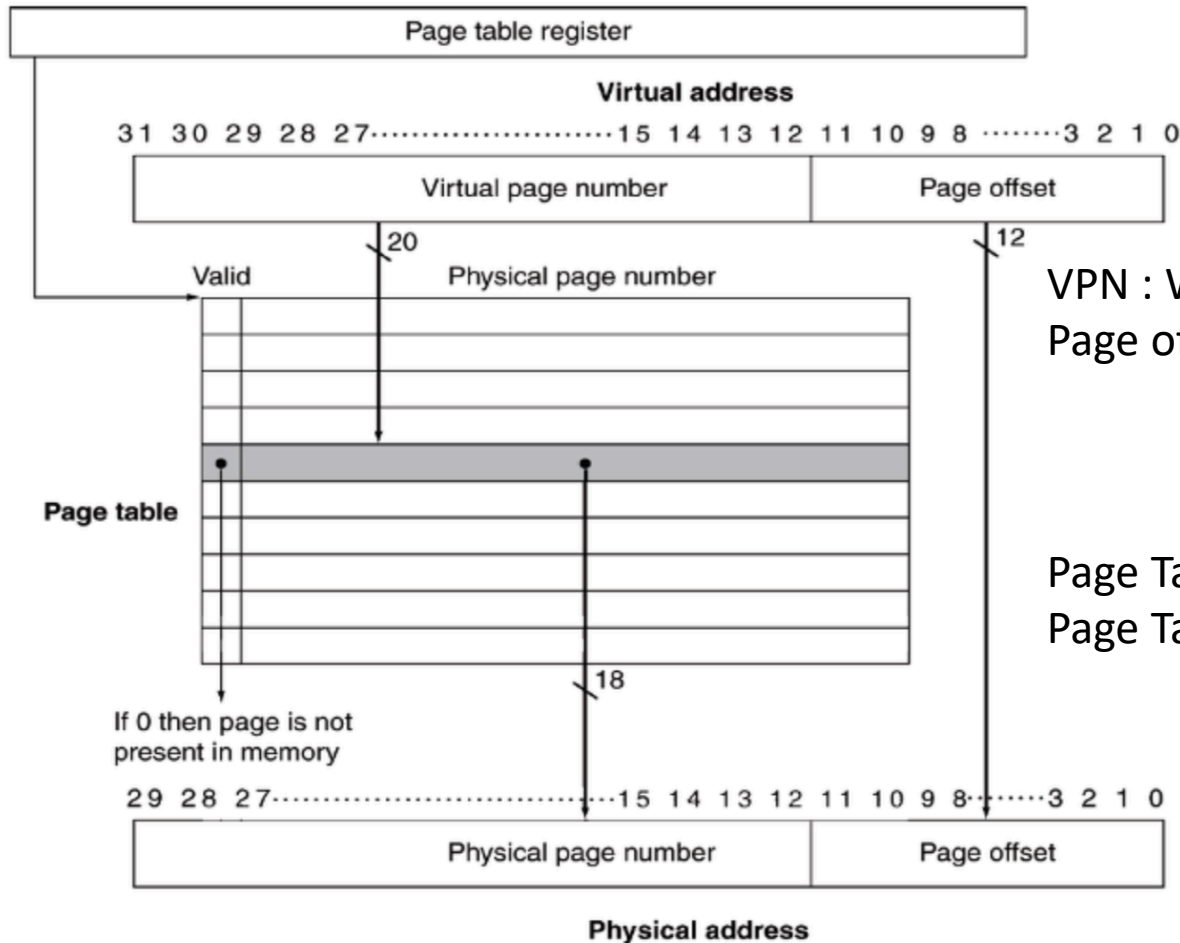
Functions calling this function: mem\_sbrk

File	Function	Line
mm-explicit.c	requestMoreSpace	172 ptrNewBlock = (void *)((unsigned int )mem_sbrk(totalSize - 4);
1 mm-explicit.c	mm_init	204 if (mem_sbrk(initsize) == (void *)-1) {
2 mm-implicit.c	mm_init	99 if ((heap_listp = mem_sbrk(4*WSIZE)) == NULL)
3 mm-implicit.c	extend_heap	227 if ((bp = mem_sbrk(size)) == (void *)-1)
4 mm-naive.c	mm_malloc	62 void *p = mem_sbrk(newsize);
5 mm-test.c	mm_init	43 mem_sbrk(64000);
6 mm-test.c	mm_malloc	54 void *p = mem_sbrk(newsize);
7 mm-tree.c	mm_init	726 if (mem_sbrk(HEAP_INITSIZE) == NULL)
8 mm-tree.c	mm_realloc	826 if (mem_sbrk(grow_size) == NULL)
9 mm-tree.c	mm_malloc	875 if (mem_sbrk(block_size) == NULL)

Find this C symbol:  
Find this global definition:  
Find functions called by this function:  
Find functions calling this function:  
Find this text string:  
Change this text string:  
Find this egrep pattern:  
Find this file:  
Find files #including this file:  
Find assignments to this symbol:



# Virtual to Physical Address Translation



VPN : Which page to look for


Page offset : Where in the page to look for

Page Table : Contains process page info

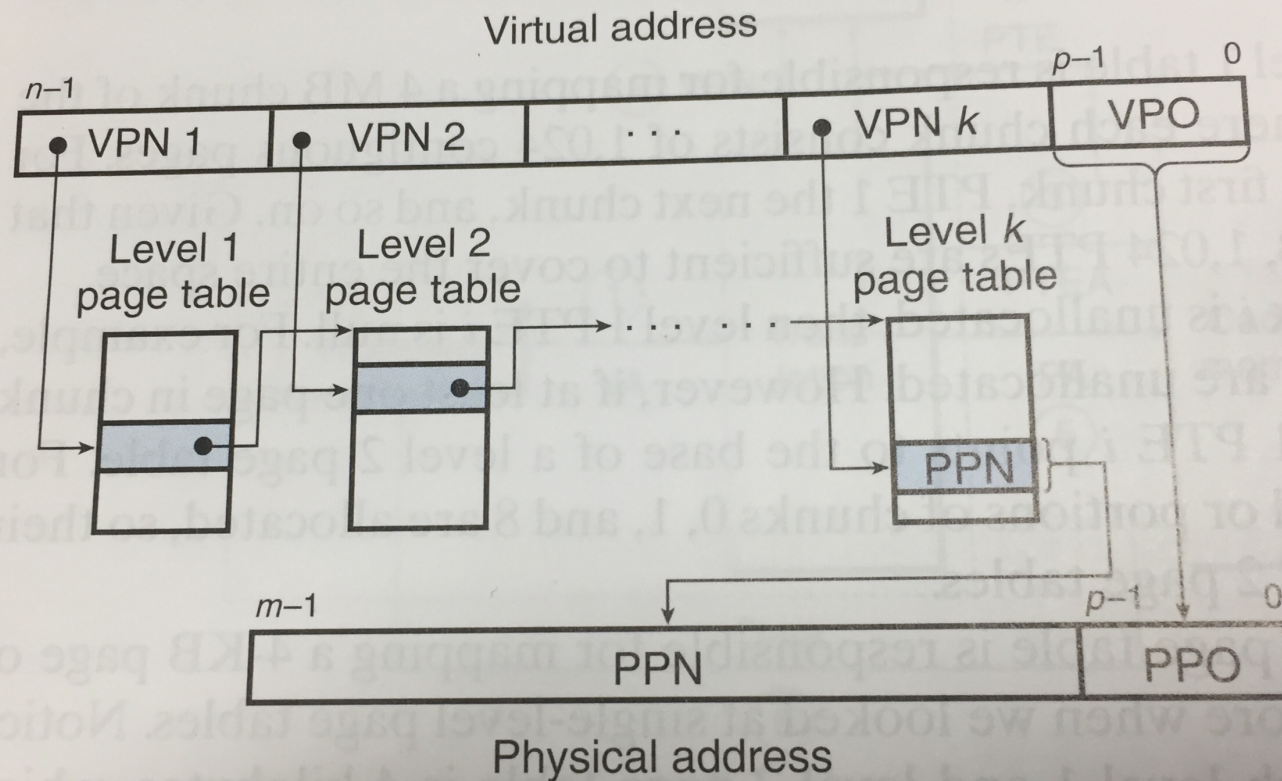
Page Table Entry : Record of Page Table

# Example

- Virtual to Physical Translation
  - 4KB page size, 20bit virtual address
  - VA : 0xF13B4



# Multilevel Page Tables



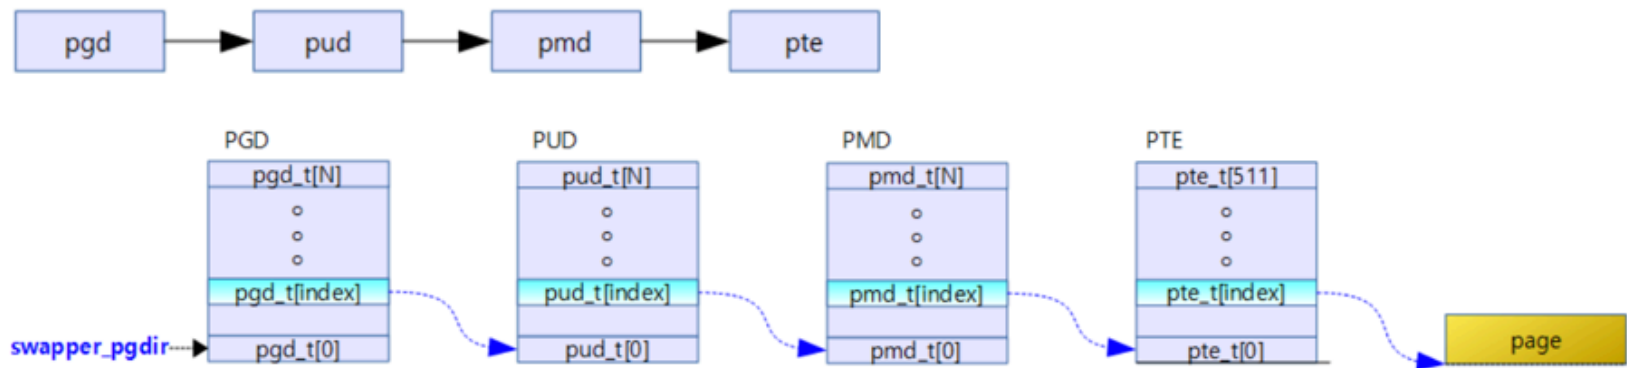
**Figure 9.18** Address translation with a  $k$ -level page table.

# Part2. Finding physical address using virtual address

- Spec
  - In app.c
    - makes a virtual address mapped to predefined physical address
  - In your kernel module
    - 1. get pid of app and virtual address
    - 2. returns physical address
  - In app.c
    - Compares return value from kernel module with predefined value.

# Hints

- Page walk API
  - `/usr/src/linux/arch/x86/include/asm/pgtable.h`
- Look for the schemes how virtual address is translated to physical address
- Page walk procedure in linux 4.4.0 (5 level in higher ver.)



# Testing your program

- Step
  - 0. `sudo su`
  - 1. `make`
  - 2. `./app`

```
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# sudo su
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# make
make -C /lib/modules/4.15.0-47-generic/build M=/home/yschoi/kernellab_full/solution/paddr modules;
make[1]: Entering directory '/usr/src/linux-headers-4.15.0-47-generic'
  CC [M]  /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.mod.o
  LD [M]  /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.15.0-47-generic'
gcc -o app app.c;
sudo insmod dbfs_paddr.ko
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# ./app
vaddr { 7ffbf301000 } paddr { 0 }
vaddr { 7ffbf301000 } paddr { 0 }
vaddr { 7ffbf301000 } paddr { 234512000 }
[TEST CASE]      PASS
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr#
```

# Asked Questions

- 0. setup
  - (1) 사전에 공지해드린 대로 delay는 받지 않습니다.
- 1. ptree
  - (1) init process의 이름이 systemd
    - -> 상관 없습니다.
  - (2) kernel message 출력 형식
    - -> 형식은 자유롭게 하셔도 됩니다.
    - 정상 동작하는 상황에서 출력되는 메시지는 제출 시 주석 처리.
    - Error case에서 출력되는 메시지는 놔두셔도 상관 없습니다.
      - -> 테스트할 때 출력되면 안되겠죠?
- 2.paddr
  - (1) huge page(2M~256M)의 case
    - -> 고려하지 않아도 됩니다.
  - (2) test input case
    - -> invalid vaddr은 test하지 않습니다.

# References

---

- Tmux guide
  - <https://tmuxguide.readthedocs.io/en/latest/tmux/tmux.html>
- Ctags
  - <https://bowbowbow.tistory.com/15>
- Cscope
  - <https://harryp.tistory.com/131>
- Address Translation, Multilevel Page table
  - P.849~855, R. E. Briant, D. R. O'Hallaron, Computer Systems, A programmer's perspective 3<sup>rd</sup> edition